

Models

Deploy and scale models on your GPU infrastructure of choice with NVIDIA NIM inference microservices

Optimized by NVIDIA

Launch from Hugging Face

Beta

Filters (1)

Free Endpoint38

Partner Endpoint14

Download Available19

Use Case

Image-to-Text5

Synthetic Data Generation4

Text Translation2

Retrieval Augmented Generation1

Text-to-Speech1

+ Show more

Inference Providers

OpenRouter10

Together AI7

Reset

Apply

GMI Cloud6

Deepinfra6

+ Show more

Publisher

NVIDIA24

Meta4

Google3

DeepSeek AI2

OpenAI1

+ Show more

NIM Container GPUs

A100 SXM4 80GB0

H100 80GB HBM30

L40S0

A100 PG509 2000

A10G0

38 models

Refine models

Q

Sort By

Most Popular

NVIDIA

Downloadable

Free Endpoint

nemotron-3-super-120b-a12b

Open, efficient hybrid Mamba-Transformer MoE with 1M context, excelling in agentic reasoning, coding, planning, tool calling, and more

moe+4

↓↑ 65M

🕒 5mo

NVIDIA

Downloadable

Free Endpoint

nemotron-3-ultra-550b-a55b

Open, efficient hybrid Mamba-Transformer MoE with 1M context, excelling in agentic reasoning, coding, planning, tool calling, and more

agent+4

↓↑ 52M

🕒 3mo

OpenAI

Downloadable

Free Endpoint

gpt-oss-20b

Smaller Mixture of Experts (MoE) text-only LLM for efficient AI reasoning and math

reasoning+3

↓↑ 19M

🕒 1y

Minimaxai

Deprecation in 6d

Free Endpoint

minimax-m3

MiniMax M3 Preview is a multimodal MoE vision-language model with strong reasoning, coding, and tool-calling capabilities.

coding+2

↓↑ 10M

🕒 2mo

NVIDIA

Downloadable

Free Endpoint

nemotron-3-nano-omni-30b-a3b-reasoning

Nemotron 3 Nano Omni is an omni-modal reasoning model that understands images, video, speech, text.

image-to-text+4

↓↑ 8M

🕒 4mo

Google

Downloadable

Free Endpoint

+ Show more

gemma-4-31b-it

Dense 31B model delivering frontier reasoning for coding, agentic workflows, and fine-tuning.

coding

+3

↓↑ 6M

🕒 5mo

Meta

Downloadable

Free Endpoint

llama-3.2-90b-vision-instruct

Cutting-edge vision-Language model exceling in high-quality reasoning from images.

image-text retrieval

+4

↓↑ 4M

🕒 1y

Google

Downloadable

Free Endpoint

diffusiongemma-26b-a4b-it

Diffusion-based 26B parameter LLM enabling parallel token generation for real-time text apps

diffusion-llm

+2

↓↑ 4M

🕒 2mo

Meta

Downloadable

Free Endpoint

llama-3.2-11b-vision-instruct

Cutting-edge vision-language model exceling in high-quality reasoning from images.

image-text retrieval

+4

↓↑ 3M

🕒 1y

NVIDIA

Downloadable

Free Endpoint

nemotron-3.5-content-safety

Multilingual, multimodal model for detecting unsafe and toxic content.

llm safety

+3

↓↑ 2M

🕒 3mo

Mistral AI

Free Endpoint

mistral-nemotron

Built for agentic workflows, this model excels in coding, instruction following, and function calling

language generation

+2

↓↑ 1M

🕒 1y

 NVIDIA

Downloadable

Free Endpoint

 Meta

Free Endpoint

llama-guard-4-12b

Multi-modal model to classify safety for input prompts as well output responses.

quantum

+3

llm multimodal safety

+3

↓↑ 442K

🕒 4mo

↓↑ 357K

🕒 1y

 NVIDIA

Free Endpoint

llama-3.1-nemotron-safety-guard-8b-v3

Leading multilingual content safety model for enhancing the safety and moderation capabilities of LLMs

content moderation

+4

↓↑ 344K

🕒 10mo

 NVIDIA

Free Endpoint

riva-translate-4b-instruct-v1_1

Translation model in 12 languages with few-shots example prompts capability.

nvidia nim

+2

↓↑ 343K

🕒 8mo

 NVIDIA

Downloadable

Free Endpoint

synthetic-video-detector

NVIDIA Synthetic Video Detector is an AI-powered micro-service for detecting AI-generated (synthetic) videos.

broadcast

+4

↓↑ 309K

🕒 4mo

 NVIDIA

Free Endpoint

magpie-tts-zeroshot

Expressive and engaging text-to-speech, generated from a short audio sample.

tts

+3

↓↑ 17K

🕒 1y

 Google

Free Endpoint

paligemma

Vision language model adept at comprehending text and visual inputs to produce informative responses

image

+8

↓↑ 12K

🕒 2y

 NVIDIA

Downloadable

Free Endpoint

Studio Voice

Enhance input speech recorded with low-quality microphones in noisy or reverberant environments, producing studio-quality speech.

broadcast

+6

↓↑ 6K

🕒 1v

 NVIDIA

Downloadable

Free Endpoint

cosmos3-nano-reasoner

Vision language model that excels in understanding the physical world using structured reasoning on videos or images.

video understanding

+8

↓↑ 2K

🕒 3mo

 NVIDIA

Downloadable

Free Endpoint

cosmos3-nano

Generates physics-aware videos from text prompts or an image prompt for physical AI development.

autonomous vehicles

+5

↓↑ 2K

🕒 3mo

 NVIDIA

Downloadable

Free Endpoint

Active Speaker Detection

Detect and track speaker identities across video frames.

broadcast

+7

↓↑ 1K

🕒 4mo

 NVIDIA

Downloadable

Free Endpoint

Background Noise Removal

Removes unwanted noises from audio improving speech intelligibility.

denoising

+3

↓↑ 1K

🕒 1y

 NVIDIA

Free Endpoint

nemotron-voicechat

Nemotron 3 Voicechat

english

+2

↓↑ 1K

🕒 5mo

 NVIDIA

Deprecation in 6d

Free Endpoint

cosmos-transfer1-7b

Generates physics-aware video world states for physical AI development using text prompts and multiple spatial control inputs derived from real-world data or simulation.

synthetic data generation

+4

↓↑ 286

🕒 1y

 NVIDIA

Free Endpoint

streampetr

StreamPETR offers efficient 3D object detection for autonomous driving by propagating sparse object queries temporally.

autonomous vehicles +3

↓↑ 179 ⌚ 9mo

 NVIDIA

Free Endpoint

bevformer

Advanced transformer for multi-frame bird's-eye-view 3D perception in autonomous driving.

autonomous vehicles +3

↓↑ 95 ⌚ 1y

 NVIDIA

Free Endpoint

sparsedrive

End-to-end autonomous driving stack integrating perception, prediction, and planning with sparse scene representations for efficiency and safety.

autonomous vehicles +3

↓↑ 74 ⌚ 1y

 NVIDIA

Free Endpoint

cosmos-transfer2.5-2b

Generates physics-aware video world states for physical AI development using text prompts and multiple spatial control inputs derived from real-world data or simulation.

synthetic data generation +4

⌚ 6mo

 DeepSeek AI

Free Endpoint

deepseek-v4-flash-0731

284B MoE (13B active) model ideal for long-context workloads optimized for coding, chat, and agentic workflows

moe +3

⌚ 15d

 DeepSeek AI

Free Endpoint

deepseek-v4-pro-0813

DeepSeek V4 scales to 1M-token context windows with efficient MoE architecture for coding tasks.

coding +3

⌚ 7d

 NVIDIA

Free Endpoint

ising-calibration-1.5-31b

NVIDIA-Ising-Calibration-1.5 is a dense multimodal vision-language model built on Gemma 4 31B. It analyzes quantum computing calibration experiment plots and...

quantum computing +3

1mo

Moonshotai

Downloadable

Free Endpoint

kimi-k3

~2.8T hybrid KDA+MLA multimodal MoE for long-horizon coding, agentic tool use, and image understanding.

multimodal +3

6d

Poolside

Free Endpoint

laguna-xs-2.1

Efficient 33B MoE for local, long-horizon agentic coding and terminal tasks

agentic ai +3

1mo

Meta

Downloadable

Free Endpoint

muse-glimmer-30b

Muse Glimmer 30B is a multimodal reasoning model accepting text and images, with native tool-calling and separate reasoning output.

multimodal +5

24d

NVIDIA

Free Endpoint

nemotron-3-embed-1b

1B embedding model for semantic search, retrieval, and RAG applications.

nemotron retriever +4

1mo

NVIDIA

Downloadable

Free Endpoint

nemotron-3.5-lightning-30b-a3b

Fastest 30B A3B MoE model with leading domain accuracy for specialized agentic tasks

customization +3

23d

NVIDIA

Free Endpoint

riva-translate-4b-instruct-v2

Translation model in 37 languages with few-shots example prompts capability.

nvidia nim +2

1mo

Items per page 96